

Shadow Models – Incremental Transformations for MPS

Markus Voelter
independent/itemis Oetztaler Strasse
38 Stuttgart 70727 Germany
voelter@acm.org

Sascha Lisson
itemis AG Industriestrasse 6 Stuttgart
70565 Germany
lisson@itemis.de

Klaus Birken, Alexander Rimer
itemis AG Industriestrasse 6 Stuttgart
70565 Germany
[birken|rimer]@itemis.de

Abstract

Shadow Models is an incremental transformation framework for MPS. The name is motivated by the realization that many analyses are easier to do on an AST whose structure is different from what the user edits. To be able to run such analyses interactively in an IDE, these “shadows” of what the user edits must be maintained in realtime. Incrementality can deliver the needed short response times. In the paper we motivate the system through example use cases, and outline the transformation framework.

CCS Concepts • **Computer systems organization** → **Embedded systems**; *Redundancy*; Robotics; • **Networks** → Network reliability;

Keywords domain-specific languages, model transformations, incrementality, language workbenches, MPS

1 Introduction

Context A core problem when representing information formally with models is that different tasks suggest different representations of the same information. For example, one particular abstract syntax might be useful for the user as she edits the model, but another representation might be more suitable for performing a particular analysis, and a third one might be suitable for execution. It is a well-known approach in any number of tools, including compilers, to transform a source model into several intermediate representations for particular kinds of analyses, and ultimately, execution.

Problem If the results of these analyses should be available to the user in an IDE, they expect the results to be available as they edit the program, in realtime. This requires that the representation of the information that is suited for the particular analysis is maintained as the user edits the program. For all but the computationally cheapest transformations and analyses, this requires an incremental maintenance (and ideally, analysis) of the derived representations: the user makes an edit to the input model, this change is propagated into the transformation engine, the target model is updated incrementally, the analysis is performed, and then the (incrementally updated) analysis results are piped back up to the user. This can potentially be done in multiple steps to form

(a pipeline), and one might also want to maintain several shadow models from a single source (fan-out).

Solution Approach Shadow models is an infrastructure for performing incremental maintenance of derived models. It is fully integrated into MPS. In this paper, we give an overview over the framework, current prototypical use cases as well as open areas for future research evolution of the tool.

2 Use Cases

2.1 Growing Domain-Specific Languages

An important approach for developing (domain-specific) languages is to grow a specialized language from a more general one, an idea beautifully illustrated for Lisp by Guy Steele’s Growing a Language [12] and Racket [6], also in Lisp. The semantics of extensions is defined through reduction (aka desugaring) to the base language.

Because of MPS’ rich support for language modularity, this approach is idiomatic. JetBrains has used this approach internally to extend Java with syntax for working with relational databases. And mbeddr extends C with domain-specific concepts for embedded software development [19]. More recently, KernelF [16] has been used as a base language for DSLs in finance and healthcare [17].

In the Java and mbeddr, the reduction to the base language happened on demand when the user invokes **Make** in the IDE. However, in addition to compilation to Java, KernelF is also executed with an in-IDE interpreter to shorten the feedback cycle and reduce the need for external build and execution infrastructure. This means that, for every DSL language construct that extends KernelF, the language engineer has to develop both a code generator to Java *and* an interpreter. This is tedious and error-prone.

Why Realtime KernelF2 is a minimal, but expressive functional language with an interpreter and a code generator. When extending KernelF2, the semantics of the new language concepts is defined through a Shadow Model-based transformation to KernelF2. Because these transformations are executed incrementally, an interpreter for the base language can now also be used for extensions. We define the semantics *once* and get a generator as well as an interactive interpreter. The overall vision, as shown in Fig. 1, also includes various verifiers as backends for the language.

2.2 Code Weaving for Safety

Separation of concerns (SoC) is proven to reduce code complexity and to support modularity and reuse. A consequence is that, for the final system, the previously separated concerns have to be reintegrated, a process known as “weaving”. In the context of the SAFE4I project¹ we use Shadow Models to

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SLE ’19, Athens

© 2019 Copyright held by the owner/author(s). Publication rights licensed to ACM. 978-1-4503-9999-9/18/06...\$15.00
DOI: 10.1145/1122445.1122456

¹<https://www.edacentrum.de/safe4i/>, BMBF FKZ 01|S17032




Figure 1. KernelF2 is a functional base language intended for extension towards DSLs. The semantics of extensions are defined via Shadow Model-based transformations. Because these are executed in incrementally, we can use interactive backends such as interpreters.

incrementally weave safety concerns into C programs written in mbeddr [18].

Separating the safety concern is feasible because most safety measures rely on a limited number of established patterns such as checksums or redundant computation with subsequent voting [8]. When the core logic and the safety patterns are kept separate, both can evolve independently and can be removed on demand. In addition, the same pattern can potentially be applied to many different target locations. It also fits well with a development process that distinguishes between safety engineers and (regular) embedded developers: each can maintain their own artifact.

Safety engineers use a DSL to specify *safety patterns* modularly. The pattern describes the constraints regarding a potential *weaving site* (in terms of structure, type system and data flow), plus the modifications to the core code (Fig. 2). The embedded software engineers mark the locations in their code where a particular safety pattern will be woven in. Finally, a *weaver*, implemented as a Shadow Model transformation merges the two concerns.

Why Realtime A drawback of SoC is that it requires re-assembling the overall system from the separated artifacts. To minimize this drawback, it is useful to show the weaving result to the user. The shorter the feedback, and the lower the requirements on the build infrastructure the better. This is especially true because some of the weavings are non-trivial; it is useful to show the result and give the safety engineer the opportunity to fix potential problems.

SAFETY-PATTERN:Safe Function – Dual Modular Redundancy	
Configuration	
Delta constraints	
<pre> constants { #alias NOTIFY_DC_DETECTION = hex«ffffff06u»; } delta function LEDStructCompare: LEDCoordinates { double fxDelta = (ORIG.fx - DUP.fx) * (ORIG.fx - DUP.fx); double fyDelta = (ORIG.fy - DUP.fy) * (ORIG.fy - DUP.fy); if (fxDelta > .5 fyDelta > .5) { *((int8 volatile*) (NOTIFY_DC_DETECTION)) = 1; } if } </pre>	
Required capabilities	

Figure 2. A safety pattern specification.

2.3 Incremental Staging of Feature Models

Feature modeling is well-established for modeling variability in product lines [9]: a *feature model* specifies the set of possible products by defining identifiable features and the constraints between them; *configurations* specify individual products by selecting features while respecting the constraints. The formalism comes with a set of predefined constraints (such as **mandatory**, **optional**, **n-of-m** and **1-of-m** but also allows custom constraints using Boolean expressions.

Sometimes the product is configured in steps, where every step makes additional selections; only the final configuration in the sequence defines a concrete product. Such *staged configuration* [2] are typically used along a supply chain or to distinguish between build-time and runtime configuration decisions.

We implemented feature models in MPS as a building block for customer-specific modeling environments. In addition to staged configuration, the tool also supports attributes, modularity via instantiation, and cardinalities [3]. The tool uses the Z3 SMT solver [4] in order to check consistency of feature models, and interactively guides the user towards valid configurations.

Creating a partial configuration C of feature model M implicitly defines [2] a specialized feature model M' by removing all features, attributes and constraints that became redundant due to the user's decisions in C ; it also leads to specialization of constraints, for example, $F_1 \Rightarrow F_2 \vee F_3$ will be specialized to $F_1 \Rightarrow F_3$, if F_2 has been deselected in C . In the next stage, a more specific configuration C' is derived from the specialized feature model M' . The creation of the derived feature models is implemented via Shadow Models. Fig. 3 shows the process.




Figure 3. Incremental staging of feature models.

Why Realtime The specialized feature model becomes available right after each user decision. This has several benefits for the user: (i) for each user decision the impact on the resulting feature model is immediately visible; (ii) the user understands at all times which downstream decisions are still open; and (iii) the solver checks on the derived feature model provide additional insights, e.g., if the specialization leads to redundant constraints.

3 Framework

3.1 The core transformation framework

The framework consists of five components: the transformation DSL, an engine for incremental computations, the transformation engine itself, an integration with MPS' model repository and various IDE integrations.

Transformation DSL The language is functional: each function takes one or more source nodes as input and produces one or more output nodes. Functions are polymorphic in all arguments and support multimethod-style dispatch [11]. The DSL exploits MPS’ strength regarding language extension and composition: queries and low-level expressions reuse MPS’ Java implementation and model access APIs. They need not be declarative, because dependency analysis happens dynamically at runtime. Reference resolution is based on (cached) re-invocation of transformation rules or explicitly defined labels; we cover this in more detail in Section 3.2. Finally, there is syntax to help with lifting analysis results from the target model back to the source(s). These are functions implemented as part of a transformation rule that attach error messages to the input of a rule when particular errors are present on the output.

Incremental Computation Engine The core engine works similar to Adapton [7]; essentially, Shadow Models map the domain of graph transformations to the general notion of incremental computations. The engine caches the result of function calls and records dependencies on other functions and mutable data for invalidation after a change. The computations are lazy: a transformation is only executed if the particular (part of the) result is accessed. This makes it suitable for IDE services where only the currently edited part of the input model is relevant for feedback to the user.

Transformation Engine The core engine expects computations to be expressed as pure functions whose results can be cached. Thus, each transformation rule expressed with the DSL is generated into a function that returns a fragment of the final output graph. Each fragment is connected to other fragments by a specification of the transformation rule and the parameter values.

The engine works on an internal data structure that is independent of MPS’ representation of syntax trees; an interface `INode` forms the API. The dynamically-maintained dependency graph is used to detect changes; a change to a dependency triggers a retransformation.

MPS Adapter The model data structure in MPS requires transactions for read and write access. The projectional editor of MPS directly writes user input to the model and updates the UI by rendering the updated model. Long running transactions, such as transformations, will block the editor’s write transaction, resulting in an unresponsive UI.

To decouple the transformations from the repository (and hence the editor), the first step in the transformation chain mirrors the MPS model into a persistent copy-on-write (COW) data structure[5] that allows reads without blocking writes. This data structure also has a mutable API with transactions, but the internal copy-on-write (COW) approach has better concurrency properties. Because the MPS projectional editor broadcasts change events anyway, maintaining this copy is computationally cheap. Fig. 4 shows the integration.

The result of the transformation can either be analyzed directly on the `INode` structure created by the engine or after materialization to an MPS AST (through another COW). The latter is slower, but has the advantage that existing MPS

analyses (such as type checks) can be used unchanged; it is also the basis for visualization in the editor.

IDE Integration Shadow Models is fully integrated into MPS. The DSL comes with editor support and type checking and is available as a language aspect (similar to the native MPS generators or type system specifications). The target models can be opened in MPS editors; editing is not possible, because this would require some form of bidirectionality, which Shadow Models do not support.

A new entry in the MPS project view, called the Shadow Repository, shows all the incrementally maintained models. Results of analyses on the target nodes can, after lifting, be annotated to the source nodes (red squiggles, markers in the gutter). Finally, there is a debugger that shows which transformation operated on which input nodes, created which outputs and ran in which forks.

3.2 References, Forks and Eagerness

Modularity MPS models are trees with cross-references (or: graphs with a single containment hierarchy). Those non-containment cross-references are particularly challenging: a reference of some type `P` between input nodes `A` and `B` must be mapped to a reference of some type `Q` between the corresponding output nodes `A'` and `B'`. To obtain `B'` from `B` in the transformation that transforms `A`, one can invoke the transformation `T` that maps `B` to `B'` again; because of caching, `B'` is not created a second time.

Labels However, for reasons of modularity, you might not want to know `T`. To achieve this, the language supports transformation labels, named mappings between nodes. The transformation `T:B->B'` would populate a label `L`, and other transformations can find `B'` knowing `B` and `L`. This way, labels are a kind of interface.

Laziness While this approach enables transformation modularity, it conflicts with laziness: to be able to retrieve `B'` from `L`, the label must already be filled; lazy computation will not work because `L` cannot know the transformation that fill it – ignorance of this dependency was the reason for labels in the first place. A static analysis might reveal the transformation, but not the (runtime) input parameters.

More generally, the model-driven engineering research roadmap by Kolovos et al. [10] identifies laziness as a core challenge in the context of graph transformations. The problem is that in a lazy system, less information is available at runtime because some parts of a transformation have not yet been executed; the issue with labels is an example. Another example is that the parent of a node in the output model





Figure 4. The transformation engine and its integration into the MPS model infrastructure.

might not yet be available when a node is accessed via a reference. If we follow the path `A.C.D.parent` in the this graph,



the node B will not yet be available because it is lazily computed when following the `A.B` child link; the transformation describes the parent-child relationship only from the parent to the child.

Forks Our solution is to compute results eagerly, but only in demarcated regions called forks. The transformations inside a fork are executed eagerly; labels can be used to look up targets, the `parent` can be retrieved. From the outside, the whole fork is lazy and when referencing nodes inside a fork from the outside, the lookup has to specify the fork. Effectively, the fork becomes part of the identify of the nodes created inside the fork.

Another consequence of the approach is that it is now possible to run a transformation multiple times, creating outputs with different identities, without adding an additional parameter to all involved transformation rules. This requirement was driven by the code weaving use case Section 2.2, where the same pattern has to be woven into target locations, and references must be resolved “locally” at each weave site.

Finally, a fork can be marked as `fixpoint`, which means that transformations are eagerly executed until no more rules apply; this requirement was driven by the KernelF2 use case (Section 2.1), which requires that extensions are reduced stepwise, until only base language concepts remain, similar to a term rewriting system [1].

Summing up, our approach does *not* solve the general problem of references and laziness; to the contrary, we revert to eager transformations. However, using forks, we limit the eagerness to well-defined scopes, and retaining the lazy nature of the overall **transformation**.

4 Related Work

For space reasons, we compare only superficially to a few related approaches. The **MPS Build Pipelines**, although using model-to-model-transformations, is not incremental. Unsuccessful experiments with running it interactively prompted the development of Shadow Models. **Incremental** transformations are not a new idea; for example, **VIATRA2** [15] supports incrementality based on the **IncQuery** [14] incremental graph pattern matching engine. **Dclare for MPS** is another incremental transformation engine that relies on constraints instead of functional transformations. A detailed comparison between the two is currently being performed. Shadow Models is not **bidirectional** [13]; it supports unidirectional transformations that maintain a trace back, as well as specific APIs to propagate analysis results back to the source. Our use cases do not require true bidirectionality, and we decided to go with the simpler specifications that come with unidirectional transformations.

5 Future Directions

Based on our experience from the initial projects described in Section 2, we have identified **a** several areas of improvement.

Scalability Incremental transformations are useful especially for large models; for small ones, rerunning transformations from scratch is feasible. Although our initial experience is promising, we will have to characterize the scalability in terms of shadow update time and memory use more thoroughly, and then identify strategies for optimization of the engine. A comparison with Dclare and IncQuery is part of this.

Scope of Change Tracking Right now, all models in the MPS workspace that use languages with Shadow Model transformations are tracked and transformed, even though the user might only be interested in a subset. Depending on circumstances, this can lead to unnecessary memory consumption. We will add a way to define a scope within which change tracking and transformation should be active.

Language Abstractions The current language exposes a several engine internals (such as forks); they are hard to understand for transformation authors. We will abstract them into concepts that are less technically motivated and easier to explain.

Improved Lifting Currently, lifting of results to the input model is expressed using generic callback functions; a more concise, more declarative syntax will be provided.

Extract the Tracking Engine The incremental computation capabilities of the core engine can be used for other purposes. In particular, we plan to implement an incremental interpreter based on the same framework. This will allow clients such as KernelF2 to not just incrementally maintain the desugared shadow model, but then also run this model incrementally (as long as it is functional), achieving a fully interactive, Excel-style reactive programming environment.

Editable Target Models We will experiment with making the target models editable in some limited way, without making the transformation bidirectional. The idea is to attach handlers to the target model that know how to make particular changes to the input.

6 Conclusions

At itemis we have had this running gag for a few years now: whenever we start to talk about some new end-user relevant feature, it takes at most 10 minutes until we end up with Shadow Models as an important part of the solution; we have several additional concrete use cases in mind beyond those described in Section 2. And as we have outlined in Section 5 there is still work to do. However, our initial experience is promising, and the we see many of the benefits of Shadow Models that we had hoped for in our running gag.

References

- [1] F. Baader and T. Nipkow. *Term rewriting and all that*. Cambridge university press, 1999.
- [2] K. Czarnecki, S. Helsen, and U. Eisenecker. Staged configuration using feature models. In R. L. Nord, editor, *Software Product Lines*, pages 266–283, Berlin, Heidelberg, 2004. Springer Berlin Heidelberg.
- [3] K. Czarnecki and C. H. P. Kim. Cardinality-based feature modeling and constraints: a progress report. In *International Workshop on Software Factories at OOPSLA’05*, San Diego, California, USA, 2005. ACM, ACM.
- [4] L. De Moura and N. Bjørner. Z3: An efficient smt solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*, pages 337–340. Springer, 2008.
- [5] J. R. Driscoll, N. Sarnak, D. D. Sleator, and R. E. Tarjan. Making data structures persistent. *Journal of computer and system sciences*, 38(1):86–124, 1989.
- [6] M. Felleisen, R. B. Findler, M. Flatt, S. Krishnamurthi, E. Barzilay, J. McCarthy, and S. Tobin-Hochstadt. A programmable programming language. *Communications of the ACM*, 61(3):62–71, 2018.
- [7] M. A. Hammer, K. Y. Phang, M. Hicks, and J. S. Foster. Adapton: Composable, demand-driven incremental computation. In *ACM SIGPLAN Notices*, volume 49, pages 156–166. ACM, 2014.
- [8] R. S. Hanmer. *Patterns for fault tolerant software*. John Wiley & Sons, 2013.
- [9] K. Kang, S. Cohen, J. Hess, W. Novak, and A. Peterson. Feature-oriented domain analysis (foda) feasibility study. Technical Report CMU/SEI-90-TR-021, Software Engineering Institute, Carnegie Mellon University, Pittsburgh, PA, 1990.
- [10] D. S. Kolovos, L. M. Rose, N. Matragkas, R. F. Paige, E. Guerra, J. S. Cuadrado, J. De Lara, I. Ráth, D. Varró, M. Tisi, et al. A research roadmap towards achieving scalability in model driven engineering. In *Proceedings of the Workshop on Scalability in Model Driven Engineering*, page 2. ACM, 2013.
- [11] T. Millstein and C. Chambers. Modular statically typed multi-methods. In *European Conference on Object-Oriented Programming*, pages 279–303. Springer, 1999.
- [12] G. L. Steele. Growing a language. *Higher-Order and Symbolic Computation*, 12(3):221–236, 1999.
- [13] P. Stevens. A landscape of bidirectional model transformations. In *International Summer School on Generative and Transformational Techniques in Software Engineering*, pages 408–424. Springer, 2007.
- [14] Z. Ujhelyi, G. Bergmann, Á. Hegedüs, Á. Horváth, B. Izsó, I. Ráth, Z. Szatmári, and D. Varró. Emf-incquery: An integrated development environment for live model queries. *Science of Computer Programming*, 98:80–99, 2015.
- [15] D. Varró, G. Bergmann, Á. Hegedüs, Á. Horváth, I. Ráth, and Z. Ujhelyi. Road to a reactive and incremental model transformation platform: three generations of the viatra framework. *Software & Systems Modeling*, 15(3):609–629, 2016.
- [16] M. Voelter. The design, evolution, and use of kernelf. In *International Conference on Theory and Practice of Model Transformations*, pages 3–55. Springer, 2018.
- [17] M. Voelter, B. Kolb, K. Birken, F. Tomassetti, P. Alff, L. Wiart, A. Wortmann, and A. Nordmann. Using language workbenches and domain-specific languages for safety-critical software development. *Software & Systems Modeling*, pages 1–24, 2018.
- [18] M. Voelter, D. Ratiu, B. Kolb, and B. Schaetz. mbeddr: instantiating a language workbench in the embedded software domain. *Automated Software Engineering*, 20(3):1–52, 2013.
- [19] M. Voelter, A. van Deursen, B. Kolb, and S. Eberle. Using c language extensions for developing embedded software: A case study. In *OOPSLA 2015*, 2015.