

Shadow Models [Tool Demo]

Incremental Transformations for MPS

Anonymous Author(s)

Abstract

Shadow Models is an incremental transformation framework for MPS. The name is motivated by the realization that many analyses are easier to do on an AST whose structure is different from what the user edits. To be able to run such analyses interactively in an IDE, these “shadows” of the user-facing AST must be maintained in realtime, and incrementality can deliver the needed short response times. Shadow Models is an incremental model transformation engine for MPS. In the paper we motivate the system through example use cases, and outline the transformation framework.

CCS Concepts •Software and its engineering →Application specific development environments; Domain specific languages;

Keywords domain-specific languages, model transformations, incrementality, language workbenches, MPS

ACM Reference format:

Anonymous Author(s). 2019. Shadow Models [Tool Demo]. In *Proceedings of Intl. Conf. on Software Language Engineering, Athens, October 21–22, 2019 (SLE ’19)*, 5 pages. DOI: 10.1145/1122445.1122456

1 Introduction

A core problem when representing information formally with models is that different tasks suggest different representations of the same information. For example, one particular abstract syntax might be useful for the user when editing the model, a second representation might be more suitable for a particular analysis, and a third one might be suitable for execution. It is a well-known approach in any number of tools, including compilers, to transform a source model into several intermediate representations for particular kinds of analyses, and ultimately, execution.

To be maximally useful, the results of analysis should be available to the user while she edits the model, interactively guiding the editing process. This requires that the representation that suits the particular analysis is maintained as the user edits the program. For all but the computationally cheapest transformations and analyses, this requires incremental maintenance (and ideally, analysis) of the derived representations: the user makes an edit to the input model,

the change is propagated the transformation engine, the target model is updated incrementally, the analysis is performed, and then the (incrementally updated) analysis results are piped back up to the user. This can potentially be done in multiple steps (to form a pipeline), and one might also want to maintain several shadow models from a single source.

Shadow models is an incremental model transformation engine, fully integrated into MPS.

In this paper, we give an overview of the framework, current prototypical use cases as well as open areas for future research and evolution of the tool.

2 Use Cases

2.1 Growing Domain-Specific Languages

An important approach for developing (domain-specific) languages is to grow a specialized language from a more general one, an idea beautifully illustrated for Lisp by Guy Steele’s Growing a Language [12] and Racket [6], also in Lisp. The semantics of extensions is defined through reduction (aka desugaring) to the base language.

Because of MPS’ rich support for language modularity, this approach is idiomatic. JetBrains has used this approach internally to extend Java with syntax for working with relational databases. And mbeddr extends C with domain-specific concepts for embedded software development [18]. More recently, KernelF [16] has been used as a base language for DSLs in finance and healthcare [X].

In case of Java and mbeddr, the reduction to the base language happened on demand when the user invokes **Make** in the IDE. However, in addition to compilation to Java, KernelF is also executed with an in-IDE interpreter to shorten the feedback cycle and reduce the need for external build and execution infrastructure. This means that, for every DSL language construct that extends KernelF, the language engineer has to develop both a code generator to Java *and* an interpreter. This duplication is tedious and error-prone.

Why Realtime KernelF2 is a minimal, but expressive functional language with an interpreter and a code generator. When extending KernelF2, the semantics of the new language concepts is defined through a *single* Shadow Model-based transformation to KernelF2. Because these transformations are executed incrementally, an interpreter for the base language can now also be used for extensions. We define the semantics *once* and get a generator as well as an interactive interpreter. The overall vision, as shown in Fig. 1, also includes various verifiers as backends for the language.

2.2 Code Weaving for Safety

Separation of concerns (SoC) is proven to reduce code complexity and to support modularity and reuse. A consequence is that, for the final system, the previously separated concerns

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
SLE ’19, Athens

© 2019 Copyright held by the owner/author(s). Publication rights licensed to ACM. 978-1-4503-9999-9/18/06...\$15.00
DOI: 10.1145/1122445.1122456




Figure 1. KernelF2 is a functional base language intended for extension towards DSLs. The semantics of extensions are defined via Shadow Model-based transformations. Because these are executed incrementally, we can use interactive backends such as interpreters.

have to be reintegrated, a process known as “weaving”. In the context of the SAFE4I project¹ we use Shadow Models to incrementally weave safety concerns into C programs written in mbeddr [17].

Separating the safety concern is feasible because most safety measures rely on a limited number of established patterns such as checksums or redundant computation with subsequent voting [8]. When the core logic and the safety patterns are kept separate, both can evolve independently and can be re woven on demand. In addition, the same pattern can potentially be applied to many different target locations. It also fits well with a development process that distinguishes between safety engineers and (regular) embedded developers: each can maintain their own artifact.

Safety engineers use a DSL to specify *safety patterns* modularly. The pattern describes the constraints regarding a potential *weaving site* (in terms of structure, type system and data flow), plus the modifications to the core code (Fig. 2). The embedded software engineers mark the locations in their code where a particular safety pattern will be woven in. Finally, a *weaver*, implemented as a Shadow Model transformation merges the two concerns.

Why Realtime A drawback of SoC is that it requires re-assembling the overall system from the separated artifacts. To minimize this drawback, it is useful to show the weaving result to the user. The shorter the feedback, and the lower the requirements on the build infrastructure the better. This is especially true because some of the weavings are non-trivial; it is useful to show the result and give the safety engineer the opportunity to fix potential problems.

2.3 Incremental Staging of Feature Models

Feature modeling is well-established for modeling variability in product lines [9]: a *feature model* specifies the set of possible products by defining identifiable features and the constraints between them; *configurations* specify individual products by selecting features while respecting the constraints. The formalism comes with a set of predefined constraints (such as **mandatory**, **optional**, **n-of-m** and **1-of-m**) but also allows custom constraints using Boolean expressions.

Sometimes the product is configured in steps, where every step makes additional selections; only the final configuration

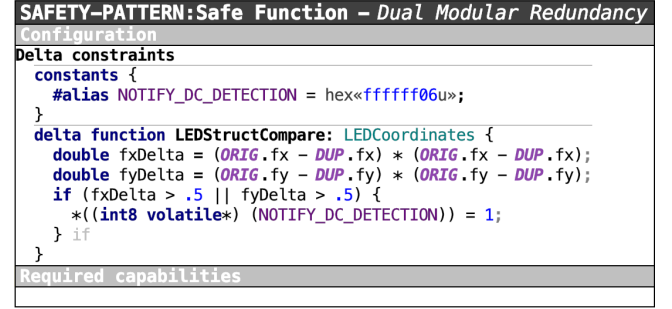


Figure 2. A safety pattern specification.

in the sequence defines a concrete product. Such *staged configurations* [2] are typically used along a supply chain or to distinguish between build-time and runtime configuration decisions.

We implemented feature models in MPS as a building block for customer-specific modeling environments. In addition to staged configuration, the tool also supports attributes, modularity via instantiation, and cardinalities [3]. The tool uses the Z3 SMT solver [4] in order to check consistency of feature models, and interactively guides the user towards valid configurations.

Creating a partial configuration C of feature model M implicitly defines [2] a specialized feature model M' by removing all features, attributes and constraints that became redundant due to the user’s decisions in C ; it also leads to specialization of constraints, for example, $F_1 \Rightarrow F_2 \vee F_3$ will be specialized to $F_1 \Rightarrow F_3$, if F_2 has been deselected in C . In the next stage, a more specific configuration C' is derived from the specialized feature model M' . The creation of the derived feature models is implemented via Shadow Models. Fig. 3 shows the process.

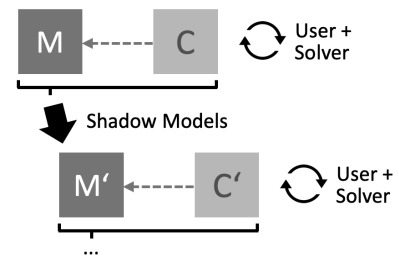


Figure 3. Incremental staging of feature models.

Why Realtime The specialized feature model becomes available right after each user decision. This has several benefits for the user: (i) for each user decision the impact on the resulting feature model is immediately visible; (ii) the user understands at all times which downstream decisions are still open; and (iii) the solver checks on the derived feature model provide additional insights, e.g., if the specialization leads to redundant constraints.

¹<https://www.edacentrum.de/safe4i/>, BMBF FKZ 01|S17032

3 Framework

3.1 The core transformation framework

The framework consists of five components: the transformation DSL, an engine for incremental computations, the transformation engine itself, an integration with MPS' model repository and various IDE integrations.

Transformation DSL The language is functional: each function takes one or more source nodes as input and produces one or more output nodes. Functions are polymorphic in all arguments and support multimethod-style dispatch [11]. The DSL exploits MPS' strength regarding language extension and composition: queries and low-level expressions reuse MPS' Java implementation and model access APIs. They need not be declarative, because dependency analysis happens dynamically at runtime. Reference resolution is based on (cached) re-invocation of transformation rules or explicitly defined labels; we cover this in more detail in Section 3.2. Finally, there is syntax to help with lifting analysis results from the target model back to the source(s). These are functions implemented as part of a transformation rule that attach error messages to the input of a rule when particular errors are present on the output.

Incremental Computation Engine The core engine works similar to Adapton [7]: the engine caches the result of function calls and records dependencies on other functions and mutable data for invalidation after a change. The computations are lazy: a transformation is only executed if the particular (part of the) result is accessed. This makes it suitable for IDE services where only the currently edited part of the input model is relevant for feedback to the user. Essentially, Shadow Models map the domain of graph transformations to the general notion of incremental computations as implemented by Adapton.

Transformation Engine The core engine expects computations to be expressed as pure functions whose results can be cached. Thus, each transformation rule expressed with the DSL is generated into a function that returns a fragment of the final output graph. Each fragment is connected to other fragments by a specification of the transformation rule and the parameter values.

The engine works on an internal data structure that is independent of MPS' representation of syntax trees. A dynamically-maintained dependency graph is used to detect changes; a change to a dependency triggers a retransformation.

MPS Adapter The model data structure in MPS requires transactions for read and write access. The projectional editor of MPS directly writes user input to the model and updates the UI by rendering the updated model. Long running transactions, such as transformations, will block the editor's write transaction, resulting in an unresponsive UI.

To decouple the transformations from the repository (and hence the editor), the first step in the transformation chain mirrors the MPS model into a persistent copy-on-write (COW) data structure [5] that allows reads without blocking writes. This data structure also has a mutable API with transactions, but the internal COW approach has better

concurrency properties. Because the MPS projectional editor broadcasts change events anyway, maintaining this copy is computationally cheap; no expensive diffs are required. Fig. 4 shows the integration.

The result of the transformation can either be analyzed directly on the `Inode` structure created by the engine or after materialization to an MPS AST (through another COW). The latter is slower, but has the advantage that existing MPS analyses (such as type checks) can be used unchanged; it is also the basis for visualization in the editor.

IDE Integration Shadow Models is fully integrated into MPS. The DSL comes with editor support and type checking and is available as a language aspect (similar to the native MPS generators or type system specifications). The target models can be opened in MPS editors; editing is not possible, because this would require some form of bidirectionality, which Shadow Models do not support.

A new entry in the MPS project view, called the Shadow Repository, shows all the incrementally maintained models. Results of analyses on the target nodes can, after lifting, be annotated to the source nodes (red squiggles, markers in the gutter). Finally, there is a debugger that shows which transformation operated on which input nodes, created which outputs and ran in which forks (explained next).

3.2 References, Forks and Eagerness

MPS models are trees with cross-references (or: graphs with a single containment hierarchy). Those non-containment cross-references are particularly challenging: a reference of some type `P` between input nodes `A` and `B` must be mapped to a reference of some type `Q` between the corresponding output nodes `A'` and `B'`. To obtain `B'` from `B` in the transformation that transforms `A`, one can invoke the transformation `T` that maps `B` to `B'` again; because of caching, `B'` is not created a second time.

Labels However, for reasons of modularity, you might not want to know `T`. To achieve this, the language supports transformation labels, named mappings between nodes. The transformation `T: B → B'` would populate a label `L`, and other transformations can find `B'` knowing `B` and `L`. This way, labels are a kind of interface.

Laziness While this approach enables transformation modularity, it conflicts with laziness: to be able to retrieve `B'` from `L`, the label must already be filled; lazy computation will not work because `L` cannot know the transformation that fills it – ignorance of this dependency was the reason for labels in the





Figure 4. The transformation engine and its integration into the MPS model infrastructure.

first place. A static analysis might reveal the transformation, but not the (runtime) input parameters.

More generally, the model-driven engineering research roadmap by Kolovos et al. [10] identifies laziness as a core challenge in the context of graph transformations. The problem is that in a lazy system, less information is available at runtime because some parts of a transformation have not yet been executed; the issue with labels is an example. Another example is that the parent of a node in the output model might not yet be available when a node is accessed via a reference. Consider this graph:



If we follow the path `A.C.D.parent`, the node `B` will not yet be available because it is lazily computed when following the `A.B` child link; the transformation describes the parent-child relationship only from the parent to the child.

Forks Our solution is to compute results eagerly, but only in demarcated regions called forks. The transformations inside a fork are executed eagerly; labels can be used to look up targets, the `parent` can be retrieved. From the outside, the whole fork is lazy and when referencing nodes inside a fork from the outside, the lookup has to specify the fork. Effectively, the fork becomes part of the identify of the nodes created inside the fork.

Another consequence of the approach is that it is now possible to run a transformation multiple times, creating outputs with different identities, without adding an additional parameter to all involved transformation rules. This requirement was driven by the code weaving use case Section 2.2, where the same pattern has to be woven into target locations, and references must be resolved “locally” at each weaving site.

Finally, a fork can be marked as `fixpoint`, which means that transformations are eagerly executed until no more rules apply; this requirement was driven by the KernelF2 use case (Section 2.1), which requires that extensions are reduced stepwise, until only base language concepts remain, similar to a term rewriting system [1].

Summing up, our approach does *not* solve the general problem of references and laziness; to the contrary, we revert to eager transformations. However, using forks, we limit the eagerness to well-defined scopes, and retaining the lazy nature of the overall transformation. Initial experience suggests that this compromise works in terms of performance and scalability, but further evaluation is necessary.

4 Related Work

For space reasons, we compare only superficially to a few related approaches. The **MPS Build Pipelines**, although using model-to-model-transformations, is not incremental. Unsuccessful experiments with running it interactively prompted the development of Shadow Models. **Incremental** transformations are not a new idea; for example, **VIATRA2** [15] supports incrementality based on the **IncQuery** [14] incremental graph pattern matching engine. **Dclare for MPS** is

another incremental transformation engine that relies on constraints instead of functional transformations. Shadow Models is not **bidirectional** [13]; it supports unidirectional transformations that maintain a trace back, as well as specific APIs to propagate analysis results back to the source. Our use cases do not require true bidirectionality, and we decided to go with the simpler specifications that come with unidirectional transformations.

5 Future Directions

Based on our experience from the initial projects described in Section 2, we have identified several areas of improvement.

Scalability Incremental transformations are useful especially for large models; for small ones, rerunning transformations from scratch is feasible. Although our initial experience is promising, we will have to characterize the scalability in terms of shadow update time and memory use more thoroughly, and then identify strategies for optimization of the engine. A comparison with Dclare and IncQuery is part of this.

Scope of Change Tracking Right now, all models in the MPS workspace that use languages with Shadow Model transformations are tracked and transformed, even though the user might only be interested in a subset. Depending on circumstances, this can lead to unnecessary memory consumption. We will add a way to define a scope within which change tracking and transformation should be active.

Language Abstractions The current language exposes several engine internals (such as forks); they are hard to understand for transformation authors. We will abstract them into concepts that are less technically motivated and easier to explain.

Improved Lifting Currently, lifting of results to the input model is expressed using generic callback functions; a more concise, more declarative syntax will be provided.

Extract the Tracking Engine The incremental computation capabilities of the core engine can be used for other purposes. In particular, we plan to implement an incremental interpreter based on the same framework. This will allow clients such as KernelF2 to not just incrementally maintain the desugared shadow model, but then also run this model incrementally (as long as it is functional), achieving a fully interactive, Excel-style reactive programming environment.

Editable Target Models We will experiment with making the target models editable in some limited way, without making the transformation bidirectional. The idea is to attach handlers to the target model that know how to make particular changes to the input.

6 Conclusions

At YYYYY we have had this running gag for a few years now: whenever we start to talk about some new end-user relevant feature, it takes at most 10 minutes until we end up with Shadow Models as an important part of the solution; we have several additional concrete use cases in mind beyond those described in Section 2. And as we have outlined in Section 5 there is still work to do. However, our initial experience is promising, and we see many of the benefits of Shadow Models that we had hoped for in our running gag.

References

- [1] F. Baader and T. Nipkow. *Term rewriting and all that*. Cambridge university press, 1999.
- [2] K. Czarnecki, S. Helsen, and U. Eisenecker. Staged configuration using feature models. In R. L. Nord, editor, *Software Product Lines*, pages 266–283, Berlin, Heidelberg, 2004. Springer Berlin Heidelberg.
- [3] K. Czarnecki and C. H. P. Kim. Cardinality-based feature modeling and constraints: a progress report. In *International Workshop on Software Factories at OOPSLA'05*, San Diego, California, USA, 2005. ACM, ACM.
- [4] L. De Moura and N. Bjørner. Z3: An efficient smt solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*, pages 337–340. Springer, 2008.
- [5] J. R. Driscoll, N. Sarnak, D. D. Sleator, and R. E. Tarjan. Making data structures persistent. *Journal of computer and system sciences*, 38(1):86–124, 1989.
- [6] M. Felleisen, R. B. Findler, M. Flatt, S. Krishnamurthi, E. Barzilay, J. McCarthy, and S. Tobin-Hochstadt. A programmable programming language. *Communications of the ACM*, 61(3):62–71, 2018.
- [7] M. A. Hammer, K. Y. Phang, M. Hicks, and J. S. Foster. Adapton: Composable, demand-driven incremental computation. In *ACM SIGPLAN Notices*, volume 49, pages 156–166. ACM, 2014.
- [8] R. S. Hanmer. *Patterns for fault tolerant software*. John Wiley & Sons, 2013.
- [9] K. Kang, S. Cohen, J. Hess, W. Novak, and A. Peterson. Feature-oriented domain analysis (foda) feasibility study. Technical Report CMU/SEI-90-TR-021, Software Engineering Institute, Carnegie Mellon University, Pittsburgh, PA, 1990.
- [10] D. S. Kolovos, L. M. Rose, N. Matragkas, R. F. Paige, E. Guerra, J. S. Cuadrado, J. De Lara, I. Ráth, D. Varró, M. Tisi, et al. A research roadmap towards achieving scalability in model driven engineering. In *Proceedings of the Workshop on Scalability in Model Driven Engineering*, page 2. ACM, 2013.
- [11] T. Millstein and C. Chambers. Modular statically typed multi-methods. In *European Conference on Object-Oriented Programming*, pages 279–303. Springer, 1999.
- [12] G. L. Steele. Growing a language. *Higher-Order and Symbolic Computation*, 12(3):221–236, 1999.
- [13] P. Stevens. A landscape of bidirectional model transformations. In *International Summer School on Generative and Transformational Techniques in Software Engineering*, pages 408–424. Springer, 2007.
- [14] Z. Ujhelyi, G. Bergmann, Á. Hegedüs, Á. Horváth, B. Izsó, I. Ráth, Z. Szatmári, and D. Varró. Emf-incquery: An integrated development environment for live model queries. *Science of Computer Programming*, 98:80–99, 2015.
- [15] D. Varró, G. Bergmann, Á. Hegedüs, Á. Horváth, I. Ráth, and Z. Ujhelyi. Road to a reactive and incremental model transformation platform: three generations of the viatra framework. *Software & Systems Modeling*, 15(3):609–629, 2016.
- [16] M. Voelter. The design, evolution, and use of kernelf. In *International Conference on Theory and Practice of Model Transformations*, pages 3–55. Springer, 2018.
- [17] M. Voelter, D. Ratiu, B. Kolb, and B. Schaetz. mbeddr: instantiating a language workbench in the embedded software domain. *Automated Software Engineering*, 20(3):1–52, 2013.
- [18] M. Voelter, A. van Deursen, B. Kolb, and S. Eberle. Using c language extensions for developing embedded software: A case study. In *OOPSLA 2015*, 2015.

Demo Outline

Use of Shadow Models for KernelF2

The demo will focus on one of the use cases, KernelF2 (discussed in Section 2.1); showing demos of all three use cases would result in too little detail for each one.

The next code snippet shows a selection of the KernelF2 base language concepts. We define a couple of values and use various operators and types. Functions, generics and algebraic data types are demonstrated as well.

```
module BaseLanguage {  
  
  val x1 = 10  
  val x2 : int = 20  
  val x3 = 10 * 20  
  val x4 = x2 > x1  
  val x5 = "Hello, " + "World"  
  val x6 = true == x4  
  val x7 = true && true  
  val x8 : int = if x2 > x1  
                  then 0  
                  else 1  
  
  fun add(i: int, j: int) = i + j  
  
  val x9 = add(30, 40)  
  
  fun<T> id(v: T) : T = v  
  val x10 : int = if id(true)  
                   then id(2)  
                   else id(1)  
  
  algebraic Maybe = | Some(int)  
                   | None()  
  
  fun extract(m: Maybe, v: int) : int = match(m)  
                                         [Some(_@n) => n  
                                          => v]  
  
}
```

Some of the expressions, such as `true && true` or `10 * 20` could be evaluated statically; and indeed, there is a Shadow Models transformation for the KernelF2 base language that performs such simplifications:

```
module BaseLanguage {  
  
  val x1 = 10  
  val x2 : int = 20  
  val x3 = 200  
  val x4 = x2 > x1  
  val x5 = "Hello, World"  
  val x6 = true == x4  
  val x7 = true  
  val x8 : int = if x2 > x1  
                  then 0  
                  else 1  
  
  fun add(i: int, j: int) = i + j  
  
  val x9 = add(30, 40)  
  
  fun<T> id(v: T) : T = v  
  val x10 : int = if id(true)  
                   then id(2)  
                   else id(1)  
  
  algebraic Maybe = | Some(int)  
                   | None()  
  
  fun extract(m: Maybe, v: int) : int = match(m)  
                                         [Some(_@n) => n  
                                          => v]  
  
}
```

Now we demonstrate language extension and desugaring. This shows three of the extension constructs defined in the `kernelf2.sugar` language: `enums`, `alt` expressions and decision tables. These are not part of the core, but defined in a modular extension; their semantics is defined through a reduction to the base language. `?maybe?` is a Boolean expression that we have added for demo purposes to avoid making up arbitrary Boolean expressions all the time.

```

module Extensions {

  enum Color { red, green, yellow, yellow }

  val aColor : Color = red

  fun decide(a: int, b: int) = alt [ a > b    => 1
                                   a == b    => 2
                                   otherwise => 3 ]

  val res =
    

|         | ?maybe? | ?maybe? | ?maybe? |
|---------|---------|---------|---------|
| ?maybe? | 1       | 3       | 5       |
| ?maybe? | 2       | 4       | 6       |


}

```

Below is the reduced version of the previous code: enums become constants that follow a particular naming convention, the `alt` expression becomes nested ifs, and the decision table is translated to nested `alt` expressions which are then in turn reduced to ifs. Notice how the reduced version contains errors that are produced by type checks of the base language: declarations must have unique names and the `<!>` expression must never show up. In this case, `<!>` is produced by the transformation because the decision table has these `?maybe?` expressions, so the type system cannot statically figure out whether the table is complete and free from overlaps. **In the demo, I will of course show how changes are propagated incrementally.**

```

module Extensions {

  val Color_red = -309736043
  val Color_green = -1313361145
  val Color_yellow = -1555955472
  val Color_yellow = -1555955472

  val aColor : int = Color_red

  fun decide(a: int, b: int) = if a > b
                                then 1
                                else if a == b then 2 else 3

  val res = if ?maybe?
            then if ?maybe?
                then 1
                else if ?maybe? then 2 else <!!>
            else if ?maybe?
                then if ?maybe?
                    then 3
                    else if ?maybe? then 4 else <!!>
                else if ?maybe?
                    then if ?maybe?
                        then 5
                        else if ?maybe? then 6 else <!!>
                else <!!>
}

```

Below we can see an error message in the source model that is lifted from the shadow. No analysis happens on the level of the enum declaration itself. Lifting errors is a core use case for Shadow Models.

```

module Extensions {

  enum Color { red, green, yellow, yellow }

  fun decide(a: int, b: int) = alt [ a > b    => 1
                                   a == b    => 2
                                   otherwise => 3 ]

  val res =
    

|         | ?maybe? | ?maybe? | ?maybe? |
|---------|---------|---------|---------|
| ?maybe? | 1       | 3       | 5       |
| ?maybe? | 2       | 4       | 6       |


}

```

[LIFTED] this is the duplicate

Implementation - Simplifications

Here is the entry point to the transformation; it contributes to a predefined transformation that attaches the results of transformations to the shadow repository, a Shadow-Models internal data structure that is also visualized in the IDE. This contribution iterates over all the models in MPS that contain a `Module` (the root concept of KernelF2, see examples above) and transforms each of them through the fork `moduleFork`. It copies its input while applying two transformations (declared elsewhere) to a fixpoint.

```
namespace Repo {  
  transformation t0 contributes to ShadowRepository.Repository [ i0: Repository ]  
  [ << ... >> ]  
  -> [ o0: Repository {  
    modules: Module {  
      name: "kf2"  
      models: map _.modules.models.where({~it => it.rootNodes.ofConcept<Module>.isNotEmpty; }) -> Model {  
        name: _.name + ".reduced"  
        rootNodes: map _.rootNodes.ofConcept<Module> -> fork moduleFork []  
      }  
    }  
  } ]  
  
  fork moduleFork [ i0: Module ] {  
    root:      copy i0  
    auto apply: Desugar.desugar  
              Simplify.simplify  
    fixpoint:  true  
  }  
}
```

The next screenshot shows the `simplify` transformation. The first entry declares an abstract transformation from `Expr` to `Expr`; the other two polymorphically override the declaration for the `LogicalNotExpr` and the `PlusExpr`. The former transforms a `!true` to `false` and the second one performs the static addition of number literals.

```
namespace Simplify {  
  abstract transformation simplify overrides ... [ alt: Expr ]  
  [ << ... >> ]  
  -> [ o0: Expr { } ]  
  
  transformation? t2 overrides simplify [ i0: LogicalNotExpr ]  
  [ i0.expr.isInstanceOf(TrueLit) ]  
  [ << ... >> ]  
  -> [ o0: FalseLit { } ]  
  
  transformation? t4 overrides simplify [ i0: PlusExpr ]  
  [ i0.left.isInstanceOf(NumLit) && i0.right.isInstanceOf(NumLit) ]  
  [ << ... >> ]  
  -> [ o0: NumLit {  
    value: i0.left.NumLit.value + i0.right.NumLit.value  
  } ]  
}
```

The concrete `simplify` transformations are defined in the same language as its abstract declaration; this is not the case for the `desugar` transformation. Only the abstract declaration is specified in the `kernelF2.core` language, with language extensions expected to provide the overrides:

```
namespace Desugar {  
  abstract transformation desugar overrides ... [ alt: Expr ]  
  [ << ... >> ]  
  -> [ o0: Expr { } ]  
}
```


Implementation - Enums

Let's take a look at the reduction of enums to constants. The respective transformation overrides the `desugar` transformation for `EnumDecls`. The transformation iterates over the literals in the `enum`, and transforms each of them into a `Constant` (note how the rule is allowed to create multiple outputs as opposed to just one, as defined in the abstract declaration). The name and the value properties of the created `Constant` are computed with Java expressions that access the properties of the input node. The mapping between the each `EnumLiteral` and the resulting `Constant` is stored in the `enumLitToConst` mapping label.

```
label enumLitToConst : EnumLit -> Constant

transformation t17 overrides Desugar.desugar [ i0: EnumDecl ]
[ << ... >> ]
-> [ o0+: map _.literals -> enumLitToConst(_) <- Constant {
    name: i0.name + "_" + _.name
    value: NumLit {
        value: (i0.name + "_" + _.name).hashCode()
    }
} ]
```

We need two more transformations. First, we have to transform every `EnumType` (as in `val x: Color = ...`) into an `IntType` (`val x: int = ...`), because all enums are transformed to `int` constants. And whenever we reference an enum literal using a `EnumLitRef` expression, we have to transform it to a reference to the particular `Constant` that has been created from the referenced literal; we can find it by querying the previously populated label.

```
transformation t19 overrides Desugar.desugar [ alt: EnumType ]
[ << ... >> ]
-> [ o0: IntType { } ]

transformation t21 overrides Desugar.desugar [ l: EnumLitRef ]
[ << ... >> ]
-> [ o0: ConstantRef {
    const -> enumLitToConst(l.lit)
} ]
```

Implementation – Error Lifting for Enums

Error Lifting relies on detecting particular error messages on output nodes of the transformation and, if one exists, back-propagating a new error to one or more of the input nodes. To this end, transformation authors override a predefined operation `liftMessage` “inside” the creator of the target node. Here is the example for the enums:

```
transformation t17 overrides Desugar.desugar [ i0: EnumDecl ]
[ << ... >> ]
-> [ o0+: map _.literals -> enumLitToConst(_) <- Constant {
    name: i0.name + "_" + _.name
    value: NumLit {..}
    op ShadowRepository.liftMessage(text: string, lifter: IMessageLifter): void {
        if (text.contains("duplicate")) {
            lifter.liftMessage("this is the duplicate", _);
            lifter.liftMessage("duplicate literal names", i0);
        }
    }
} ]
```

The operation implementation, written inside the `Constant`, checks if the error message(s) reported by MPS on this `Constant` contain the word “*duplicate*”. If so, we propagate a message “*this is the duplicate*” to the currently transformed literal, and another error “*duplicate literal names*” to the input `EnumDecl`. The framework takes care automatically of removing the lifted errors if their causes go away.

Implementation – Alt Expressions

We conclude the demo with the transformation of the `AltExpression`. What makes this interesting is that a flat list of alternatives must be transformed to a tree of nested `IfExpressions`. Transforming lists to trees is achieved conveniently with the fold function in functional programming, which is why we have added a corresponding transformation primitive to the Shadow Model transformation language.

A `foldR` function joins output for the n -th element of the list to what has been constructed for the $n-1$ -th list element. In this case, we create a new `if` that reuses the condition of the current option and puts its value into the `then` part. The `else` part of the currently constructed `if` is whatever the previous iteration has created, represented by the `acc` (short for accumulator) expression inside the `foldR`. For the first entry in the `alt`’s list of options, we pass a seed value to `foldR`; in our case it is the `NeverLit`, rendered as `<!>`.

```
transformation t0 overrides Desugar.desugar [alt: AltExpr]
[ << ... >> ]
-> [ o0: foldR alt.cases, NeverLit {
    op ShadowRepository.liftMessage(text: string, lifter: IMessageLifter): void {
        lifter.liftMessage("This alt is not guaranteed to succeed; unhandled combination.");
    }
  }, IfExpr {
    cond: copy it.cond
    thenPart: copy it.val
    elsePart: ElsePart {
        expr: acc
    }
  }
]
```

This transformation will create a set of nested `ifs`, where the last one always has an `<!>` in its `else` part. However, if we look above, the transformation of the `alt`-expression in the `Extensions` module did not include an `else <!>` at the tail. Why is this? The reason is the `otherwise` in the last option, it’s basically a catch-all clause. That last option will be transformed to by the transformation above:

```
if true then 3 else <!>
```

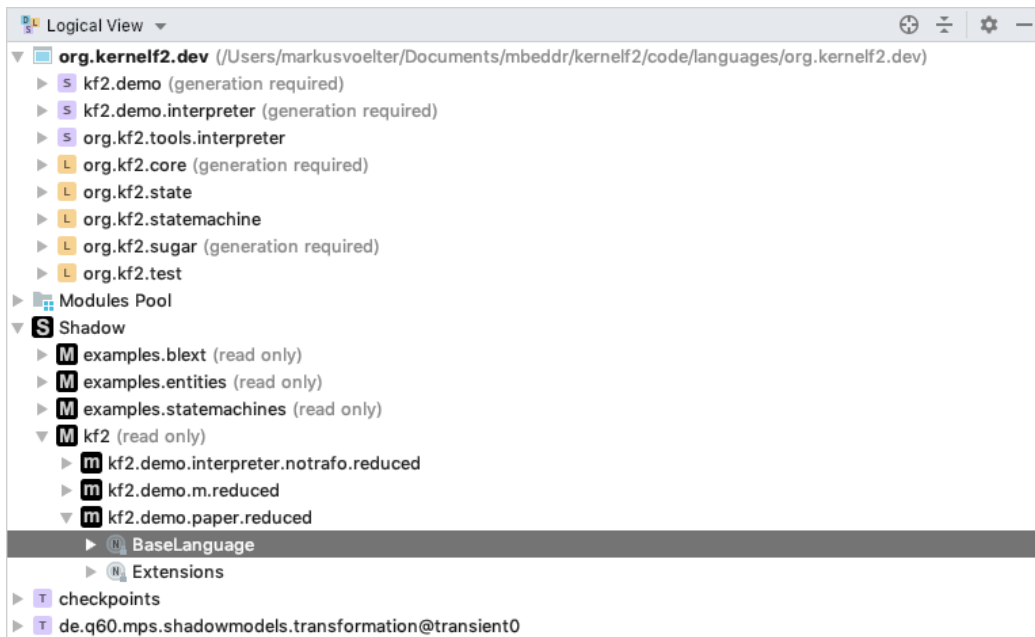
However, the set of base language simplifications defines one that transforms an `if true` to the `then` part, because we know statically the `else` option will never occur:

```
transformation? t39 overrides simplify [i0: IfExpr]
[ i0.cond.isInstanceOf(TrueLit) ]
[ << ... >> ]
-> [ o0: copy i0.thenPart ]
```

The case above thus simply becomes 3 and the `else <!>` goes away. As a corollary, this means that whenever a `<!>` is created as the result of the transformation of an `alt` expression (and survives simplification), then this is an indication of an error in the `alt` expression; which is why we add a corresponding lifter to the transformation, as can be seen above.

IDE Integrations

The Shadow Repository in the MPS project view that shows the results of all transformations; double clicking any root opens it in an editor, and the incremental transformations can be observed in realtime.



The Fork Explorer shows the execution of the transformations, including the inputs it processes and the output nodes it creates.

